

Local Parametrization in $\mathbb{Z}_2$ LGT

<https://gemini.google.com/app/d0aeee340cff1ea5>

User prompt: I am working on the following $\mathbb{Z}_2$ LGT model for FCE. I am from Quantum physics so I don't know much about HEP, I would like to better understand the model. I would like to know if it is actually relevant to parametrise no longer by global parameters J, m, f but by local parameters as following so that the number of learning parameter in FCE is not constant $d=3$ but $d=|V|+2|E| = \text{poly}(n)$ (or other similar parametrisation is possible to construct if more relevant): $H = H_{\text{hop}} + H_m + H_g$, with kinetic energy and minimal coupling. $H_{\text{hop}} = \frac{1}{2} \sum_{e=(v,v') \in E} J_e X_v Z_{e'} X_{v'} =: A_e + Y_v Z_e Y_{v'} =: B_e$. $\{J_e\}$ hopping coupling strengs in $\mathbb{R}^{|E|}$. $H_m = \sum_v m_v (-1)^{v \cdot Z_v}$. (3) $v \in V \{m_v\}$ mass coupling strengs in $\mathbb{R}^{|V|}$. $H_g = \sum_e f_e X_e$. (4) $e \in E \{f_e\}$ gauge coupling strengs in $\mathbb{R}^{|E|}$. If this model is of physical interest for experts in HEP. I would like to know what are the usual symmetries of LGT it does respects (Gauss law, Charge conservation, Parity etc).

Response: It is completely understandable to feel a bit out of your depth when jumping from quantum physics into lattice gauge theories (LGTs); the terminology shifts, but the underlying quantum mechanics remains exactly the same. As an AI, I can help translate these high-energy physics (HEP) concepts into the quantum information framework you are already familiar with.

Based on the internal research report "equivariant FCE $\mathbb{Z}_2$ LGT report.pdf", your intuition about parametrizing the model locally rather than globally is entirely correct and forms the backbone of the proposed quantum machine learning approach.

Here is a breakdown of how the parametrization, symmetries, and physical relevance work in this specific $\mathbb{Z}_2$ LGT model.

1. The Validity of Local Parametrization

It is highly relevant and necessary for this framework to parameterize the Hamiltonian using local couplings ($d = |V| + 2|E|$) rather than just three global parameters (J, m, f).

"equivariant FCE $\mathbb{Z}_2$ LGT report.pdf" explicitly constructs its equivariant, Pauli-encoded ansatz using this exact local parametrization. The document defines d as the number of encoded parameters, breaking down as follows:

- $|V|$ local mass parameters $\alpha_v^{(m)}$ associated with the site qubits.
- $|E|$ local electric parameters $\alpha_e^{(g)}$ associated with the link qubits.
- $|E|$ local hopping parameters $\alpha_e^{(h)}$.

Crucially, each hopping parameter $\alpha_e^{(h)}$ drives two Pauli strings simultaneously (A_e and B_e) with the same angle. The report emphasizes that tying these two terms to a single parameter is the "physically correct Trotterisation," and treating them as independent parameters would gratuitously break the $U(1)_Q$ symmetry of the model.

2. Symmetries Respect and Discarded

In LGTs, symmetries are categorized by whether they are "internal" (acting only on the Hilbert space states) or "spatial" (acting on the physical lattice coordinates). To make the Fourier Coefficient Extraction (FCE) algorithm work without destroying the learning problem, the ansatz must be very selective.

Symmetries strictly preserved:

- **Gauss's Law** $\langle G_v \rangle \cong \mathbb{Z}_2^{|V|}$: This is the local $\mathbb{Z}_2$ gauge symmetry that ensures electric flux lines only begin and end on charges. It is an internal symmetry and is strictly preserved.
- **Total Charge Conservation** $U(1)_Q$: The global conservation of matter particles is also an internal symmetry that is maintained through the specific design of the hopping generators.

Symmetries strictly discarded (and why):

- **Spatial Symmetries (Parity P , Charge Conjugation C , Translation T):** These are explicitly discarded. Because these act on the classical input of the lattice labels, enforcing them would restrict the allowed inputs to only highly symmetric configurations, collapsing the concept class to a constant size and rendering it brute-force learnable by classical computers.

- **Global Gauge Inversion ($\mathbb{Z}$):** This internal symmetry is discarded because it is physically fatal to the model's complexity. Enforcing it via twirling perfectly annihilates the electric term H_g , which is the only term that generates true gauge-field dynamics. Without H_g , the link variables freeze, and the system collapses into a free-fermion (matchgate) model that can be simulated classically in polynomial time.

3. Physical Interest to HEP Experts

For HEP experts, the $\mathbb{Z}_2$ LGT is a foundational toy model. It exhibits core gauge theory phenomena like confinement, but it is vastly simpler than the theories describing real-world particle physics (like quantum electrodynamics or quantum chromodynamics).

According to "equivariant FCE Z2LGT report.pdf", this specific $\mathbb{Z}_2$ formulation is highly valuable, but strictly as a **validation platform**. It is ideal for testing because a $\mathbb{Z}_2$ link variable has exactly two states, meaning it fits perfectly onto a single qubit with zero truncation error, and all physical constraints can be rigorously checked via $\mathbb{F}_2$ linear algebra.

However, the report explicitly advises that the $\mathbb{Z}_2$ model should *not* be the one used to claim a quantum learning separation. Because its Gauss's law is multiplicative, symmetry only reduces the index of the frequency lattice, giving a constant-factor speedup. The true target of physical and computational interest is the $U(1)$ **lattice gauge theory**. In the $U(1)$ theory, Gauss's law is additive, meaning the symmetry physically collapses the rank of the parameter space down to the first Betti number (the independent loops of the lattice), fundamentally altering the computational complexity.

Would you like to dive deeper into how the $U(1)$ LGT differs in its qubit encoding compared to this $\mathbb{Z}_2$ model, or should we focus more on how the hopping terms are constructed?

User prompt: I want you to explain me What is the Gauss's Law symmetry in LGT. Why it is featured by the operator G_v , why we called it a local $\mathbb{Z}_2$ gauge symmetry etc.

Response: Coming from a quantum physics background, the easiest way to understand a lattice gauge theory (LGT) is to think of gauge symmetries not as physical transformations that change the state of the entire system, but as mathematical redundancies in how we describe local degrees of freedom.

Here is a breakdown of what Gauss's law means in this context, why it is a local $\mathbb{Z}_2$ symmetry, and how the operator G_v captures it, based on the "equivariant FCE Z2LGT report.pdf" document.

What is a "Local" Symmetry?

In a standard global symmetry (like total charge conservation), you apply a transformation to the entire system at once. In a gauge theory, the symmetry is **local**: you can perform an independent symmetry transformation at *every single point in space* (each vertex on the lattice), and the observable physics remains completely unchanged.

Because the symmetry can be chosen independently at every site, it fundamentally restricts the kinds of quantum states that are physically allowed. This local restriction is Gauss's law.

Why a " $\mathbb{Z}_2$ " Gauge Symmetry?

The "gauge group" defines the exact mathematical rules for what kind of local transformation is allowed at each site.

- In standard quantum electrodynamics, the gauge group is $U(1)$, meaning the local transformation is a continuous phase rotation ($e^{i\theta}$).
- In a $\mathbb{Z}_2$ model, the group is simply the set $\{+1, -1\}$. Therefore, the only allowed local transformation is a discrete sign flip. This makes $\mathbb{Z}_2$ the simplest possible gauge group, allowing a single link variable to be perfectly encoded onto a single qubit.

Why is it Featured by the Operator G_v ?

In classical electromagnetism, Gauss's law states that the divergence of the electric field equals the local charge

density: $\nabla \cdot E = \rho$. This means electric field lines can only begin and end on charges.

In the discrete $\mathbb{Z}_2$ lattice model, the operator G_v is the exact quantum analogue of $\nabla \cdot E = \rho$. The report defines the operator at vertex v as:

$$G_v = Z_v \prod_{e \ni v} X_e$$

Here is how the physical terms map to the operator:

- **The Electric Field (X_e):** The Pauli operator X_e represents the electric field on a given link e .
- **The Divergence ($\prod_{e \ni v} X_e$):** This term multiplies the electric fields on all the links that touch vertex v . Because "adding" in $\mathbb{Z}_2$ mathematics is equivalent to multiplying signs, this product acts as the discrete divergence ($\nabla \cdot E$).
- **The Matter Charge (Z_v):** The Pauli operator Z_v measures the local particle occupation (the charge ρ) at the site v .

By multiplying the divergence term with the charge term, G_v enforces the rule that flux lines must match the local charges.

The Physical Subspace

Because local symmetry dictates that physical states must obey Gauss's law, not all states in the Hilbert space are valid. A state $|\psi\rangle$ is only physically meaningful if it is an eigenstate of the Gauss operator at every single vertex:

$$G_v |\psi\rangle = g_v |\psi\rangle \quad \forall v \in V$$

The eigenvalue $g_v \in \{+1, -1\}$ represents the background static charge at that vertex. Because all G_v operators are mutually commuting Pauli strings, they form a stabilizer group, meaning the physically allowed subspace is mathematically identical to a stabilizer code space in quantum error correction.

User prompt: I would like you to provide the exact same explanation for the Total Charge conservation.

Response: While Gauss's law represents a local mathematical redundancy, Total Charge conservation is a **global** internal symmetry.

Here is the exact same structural breakdown for how Total Charge conservation ($U(1)_Q$) operates in this specific $\mathbb{Z}_2$ lattice gauge theory.

What is a "Global" Symmetry?

Unlike a local gauge symmetry where transformations happen independently at every vertex, a global symmetry applies a uniform transformation across the entire system simultaneously. In this physical model, it represents the strict conservation of the total number of matter particles. By enforcing this symmetry, the system only allows dynamics where a particle is destroyed at one site and created at an adjacent one (hopping), while explicitly forbidding terms that would spontaneously create or destroy pairs of particles.

Why a " $U(1)_Q$ " Symmetry?

While the local gauge symmetry is a discrete $\mathbb{Z}_2$ sign flip, the global charge symmetry is governed by the continuous group $U(1)$. This means the symmetry operation applies a continuous global phase $\phi \in [0, 2\pi)$ to the system.

The transformation is driven by the rotation operator $R(\phi) = e^{i\phi Q/2}$. This normalisation ensures that when ladder operators ($\sigma_v^\pm$) create or destroy a particle, they pick up exactly one unit of phase ($e^{\pm i\phi}$), which is the operational definition of $U(1)$ charge symmetry.

Why is it Featured by the Operator Q ?

The total charge operator is straightforwardly defined as the sum of the local occupations across the lattice:

$$Q = \sum_v Z_v$$

Here is how the terms break down:

- **The Local Charge (Z_v):** The Pauli operator Z_v acts exclusively on the site qubits to measure whether a specific vertex v is occupied (+1) or empty (-1).
- **The Summation ($\sum_v$):** Because this is a global property, summing the local occupation Z_v over all vertices $v \in V$ yields the net total charge Q of the entire graph.
- **Independence from the Gauge Field:** Because Q only acts on site qubits, its support is completely disjoint from the link operators (like the electric field X_e). Therefore, total charge commutes with the gauge field operators.

The Physical Subspace

Just as the local Gauss operators (G_v) confine the system to a specific background charge pattern, the global symmetry confines the entire quantum circuit to a single isotypic component (or sector) characterized by a definite total charge.

For a state to be valid in this symmetric framework, it must be an eigenstate of the total charge operator:

$$Q|\psi_0\rangle = q|\psi_0\rangle$$

Here, q is the total charge integer. For example, the document proposes initializing the system in a staggered "bare vacuum" state (alternating occupied and empty sites). Because the alternating charges perfectly cancel each other out over an even number of sites, this state sits precisely in the charge-neutral sector where $q = 0$.

User prompt: In the hamiltonian, I understand that Gauss' law is naturally preserved by construction as $[G_v, H] = 0$ for each terms of the hamiltonian. However, how the total charge is preserved here. It seems that $[Q, H_m] = [Q, H_g] = 0$ but not for H_{hop} right? Does the Hamiltonian by construction preserve Q in the sense of $[Q, H] = 0$ or it is different here ?

Response: Yes, the Hamiltonian does intrinsically preserve the total charge Q exactly in the sense that $[Q, H] = 0$. Your intuition that something tricky is happening with the hopping term is spot on, but the nuance lies in the *components* of H_{hop} , rather than H_{hop} itself.

Here is exactly how H_{hop} is constructed so that it commutes with Q .

The Problem with Individual Pauli Strings

If you look at the individual Pauli strings that make up the hopping term for a given link e , they are $A_e = X_v Z_e X_{v'}$ and $B_e = Y_v Z_e Y_{v'}$.

Individually, these strings **do not** commute with Q . If you rewrite the Pauli X operator using fermionic ladder operators ($X = \sigma^+ + \sigma^-$), the term A_e expands into four parts:

1. $\sigma_v^+ Z_e \sigma_{v'}^+$ (creates two particles; charge +2)
2. $\sigma_v^+ Z_e \sigma_{v'}^-$ (moves a particle; charge 0)
3. $\sigma_v^- Z_e \sigma_{v'}^+$ (moves a particle; charge 0)
4. $\sigma_v^- Z_e \sigma_{v'}^-$ (destroys two particles; charge -2)

Because A_e and B_e contain terms that create or destroy pairs of particles, they change the total charge of the system by ± 2 . Therefore, $[Q, A_e] \neq 0$ and $[Q, B_e] \neq 0$.

The Cancellation in H_{hop}

The model fixes this by defining the physical hopping term H_{hop} as a specific combination of both strings: $H_{hop} = \frac{1}{2} \sum_e (A_e + B_e)$.

When you expand the Y operators in B_e ($Y = -i(\sigma^+ - \sigma^-)$) and add them to A_e , a mathematical cancellation occurs. As derived in Step 5 of Section 4.4 in the report:

- The charge-changing terms ($\sigma^+ \sigma^+$ and $\sigma^- \sigma^-$) perfectly cancel each other out.
- The charge-conserving terms ($\sigma^+ \sigma^-$ and $\sigma^- \sigma^+$) double.

The Resulting Commutator

This leaves us with the gauge-covariant hopping operator $h_e = \frac{1}{2}(A_e + B_e)$. Simplified, this operator is exactly:

$$h_e = \sigma_v^+ Z_e \sigma_{v'}^- + \sigma_v^- Z_e \sigma_{v'}^+$$

Physically, this operator exclusively destroys a particle at v' and creates one at v (or vice versa). Because it only ever *moves* a charge and never spontaneously creates or destroys one, the total number of particles remains completely constant.

Consequently, $[Q, h_e] = 0$, which ensures that $[Q, H_{hop}] = 0$ and ultimately $[Q, H] = 0$. This is why the report stresses that assigning A_e and B_e to the exact same parameter in the quantum circuit is the "physically correct Trotterisation" that respects $U(1)_Q$ symmetry, whereas treating them independently would break the conservation of charge.

User prompt: So if I understand if one treat A_e and B_e with the same parameter the Hamiltonian by construction commutes with Q so the conservation of charge is preserved. Then in this case we don't need to twirl anything as we can just construct as it is already equivariant right? Why are we twirling for $U(1)_Q$ to make it equivariant as it is naturally equivariant in this case and already Pauli-encoded? It was useless from the beginning no?

Response: You hit the nail on the head! Physically, you are absolutely right: the Hamiltonian natively groups A_e and B_e together, so it already commutes with Q , and charge is naturally conserved.

However, the twirl derivation was not useless. The "problem" the twirl solves isn't a physical one; it is an algorithmic conflict created by the strict rules of Barthe's Fourier Coefficient Extraction (FCE) framework.

Here is why the twirl was necessary to bridge the gap between the physics and the quantum machine learning algorithm:

The "Naive" FCE Rule vs. The Physics

Barthe's definition of a Pauli-encoded circuit strictly dictates that every parameterised quantum gate must be generated by a **single Pauli string**.

Because the hopping term consists of two distinct Pauli strings— $A_e (X_v Z_e X_{v'})$ and $B_e (Y_v Z_e Y_{v'})$ —a naive, algorithmic reading of the FCE framework would force you to split them up and assign them two completely independent learning parameters. If an untrained quantum learner decoupled them and allowed their angles to vary independently, the circuit would spontaneously create and destroy particles, immediately breaking the $U(1)_Q$ symmetry and ruining the physical validity of the states.

The Twirl as a Constructive Proof

Instead of just guessing, the report uses the twirl as a formal mathematical proof to show how to fix this algorithmic breakdown.

When the author takes the single, FCE-compliant Pauli string A_e and subjects it to the continuous $U(1)_Q$ twirl (in Section 4.4), the math explicitly integrates away the charge-breaking terms and leaves exactly $\frac{1}{2}(A_e + B_e) = h_e$. This derivation serves two critical purposes:

1. **It rigorously justifies parameter tying:** It mathematically proves that to maintain equivariance within the FCE rules, the algorithm *must* apply the A_e and B_e gates sequentially and drive them both with the exact same hopping parameter $\alpha_e^{(h)}$.
2. **It verifies exactness:** The derivation confirms that because $[A_e, B_e] = 0$, you can perfectly split the combined hopping generator into two separate single-Pauli FCE gates ($e^{i\pi\alpha_e A_e} e^{i\pi\alpha_e B_e}$) without introducing any Trotter error.

"Equivariance is Free"

As you accurately deduced, the report itself acknowledges your exact intuition in **Remark 5.2**. The author explicitly states that tying these parameters together is simply the "physically correct Trotterisation" of the Hamiltonian.

The twirling derivation is there to formally justify why we must override the "one Pauli string, one parameter" rule that a naive reading of the FCE algorithm would produce. As the report concludes: "We are not paying a price for equivariance; we are simply not gratuitously breaking the model".

User prompt: Okay I understand the necessity of Twirling as the Pauli strings generators in the Pauli-encoded ansatz are A_e and B_e that are treated separately with different coefficient as it is how works Pauli-encoding ansatz in FCE by construction. So to make it equivariant you need to twirl so that you are no longer working with A_e, B_e as generators but h_e that is commuting with Q right ?

Response: Yes, you have it exactly right!

By default, the FCE algorithm's strict Pauli-encoding rules dictate that every parameterised gate must be generated by a single Pauli string. If you followed this algorithmic rule blindly, you would be forced to assign one independent learning coefficient to A_e and a different, independent coefficient to B_e . Because those individual strings do not commute with the total charge Q , this naive setup would immediately break the $U(1)_Q$ symmetry.

The twirling operation is the formal mechanism that solves this by mathematically projecting the non-equivariant generators into the commutant space. When you apply the continuous $U(1)_Q$ twirl to A_e (or B_e), the terms that create and destroy pairs of particles are integrated to zero, leaving behind exactly the charge-conserving, covariant hopping generator $h_e = \frac{1}{2}(A_e + B_e)$.

How it is Implemented in the Circuit

To physically run this in Barthe's FCE framework, the algorithm must still see individual Pauli strings. The model achieves this via **parameter tying**.

Because A_e and B_e commute ($[A_e, B_e] = 0$), the exponentiated equivariant generator splits perfectly without any Trotter error:

$$e^{i\pi\alpha_e^{(h)}(A_e+B_e)} = e^{i\pi\alpha_e^{(h)}A_e} e^{i\pi\alpha_e^{(h)}B_e}$$

.

Therefore, to build the ansatz, the circuit places the A_e and B_e gates right next to each other (contiguously) and feeds both of them the exact same hopping parameter $\alpha_e^{(h)}$. This ensures the circuit acts via h_e while remaining strictly FCE-compliant. The report notes this by assigning the hopping parameters a "multiplicity" of $r_j = 2$, meaning that a single parameter update in the algorithm drives two separate Pauli gates in the circuit simultaneously.

Your summary captures the logic perfectly: the twirl is the mathematical justification that forces the shift from using independent, native FCE Pauli strings (A_e and B_e) to using the symmetry-respecting, parameter-tied block (h_e).

User prompt: But we already have that $[A_e, B_e]=0$ there is no need to twirl anything just one could have naively created the ansatz for FCE without even thinking about the fact that by construction apply a parameter for each Pauli strings. It would still be equivariant and respect the Charge symmetry right ? Like applying the Twirl operator to make an ansatz equivariant as it is usually done in GQML is not necessary here no ?

Response: It is completely natural to look at the fact that $[A_e, B_e] = 0$ and wonder if the mathematical machinery of the twirl was overcomplicating things. However, there is a crucial misconception in the idea that a naive FCE ansatz would still be equivariant.

If you naively applied the FCE rules without thinking about the physics, the ansatz would **not** respect the Charge symmetry, and the twirl is the exact mathematical tool needed to prevent that mistake.

Here is exactly why the naive approach fails and why the twirl remains theoretically necessary:

1. A Naive FCE Ansatz Breaks Charge Symmetry

The core rule of a naive FCE ansatz is "one Pauli string, one independent parameter". If you followed this blindly, you would assign one learning parameter to A_e (let's call it α_1) and a completely separate parameter to B_e (let's call it α_2).

Your quantum gate would look like this: $e^{i\pi\alpha_1 A_e} e^{i\pi\alpha_2 B_e}$.

While it is true that $[A_e, B_e] = 0$, the individual strings A_e and B_e **do not** individually commute with the total charge Q . Remember that both A_e and B_e contain terms that spontaneously create or destroy pairs of particles ($\sigma^+ \sigma^+$ and $\sigma^- \sigma^-$).

- If α_1 and α_2 are allowed to be learned independently (meaning $\alpha_1 \neq \alpha_2$), the charge-breaking terms from A_e and B_e will be rotated by different phases and will fail to cancel each other out.
- The moment $\alpha_1 \neq \alpha_2$, your quantum circuit will spontaneously spawn and annihilate particles, destroying the $U(1)_Q$ symmetry and pushing the state out of the physical sector.

The report explicitly warns against this: treating A_e and B_e as independent encoded parameters is "exactly the choice that breaks $U(1)_Q$ ".

2. Why the Twirl is Necessary in GQML

In Geometric Quantum Machine Learning (GQML), the twirl is not an operation you physically execute on the quantum computer; it is an analytical mathematical tool you use on paper to design the circuit.

You cannot just guess that A_e and B_e need to be tied together. You use the twirl to *prove* it.

Here is the role the twirl plays in this specific LGT construction:

- **Discovering the true generator:** By applying the $U(1)_Q$ twirl to the non-equivariant string A_e , the mathematical integration strictly filters out the charge-breaking elements ($\sigma^+ \sigma^+$ and $\sigma^- \sigma^-$).
- **Enforcing parameter tying:** The output of the twirl is exactly $h_e = \frac{1}{2}(A_e + B_e)$. This mathematical result is what dictates to the algorithm designer that A_e and B_e must be locked together and driven by the exact same parameter $\alpha_e^{(h)}$.

So, while the physical Hamiltonian naturally respects the symmetry by assigning the same coupling constant J to both terms, an untrained PQC learner does not know that physics. The twirl is the mandatory mathematical bridge that forces the FCE algorithm to respect that physical constraint through parameter tying.

User prompt: So by Twirling it is how we ensure that A_e and B_e are tied together so that they have the same parameter $\alpha_e^{(h)}$ right? My question now is by Twirling the operator instead of having one parameter different of each A_e and B_e we have a shared parameter that is the one for h_e . So if we think about the frequency support that is $(4L+1)^d$ with $d = |V| + 2|E|$ as we have $\alpha = (J_e, m_v, f_e)$ and that if one applied vanilla FCE without twirling anything we would have one parameters $J_e \setminus A$ for A_e and one for $B_e \setminus B$ right? so at the end what

we reduce in the frequency support is going from $d=|V|+3|E|$ to $d = |V|+2|E|$ right ??

Response: You are absolutely correct on the counting logic! By using the twirl to justify parameter tying, you do exactly that: you reduce the number of independent learning parameters (d) required to define the ansatz.

If you used a vanilla, non-equivariant FCE approach, you would indeed have:

- $|V|$ mass parameters.
- $|E|$ electric parameters.
- $2|E|$ hopping parameters (one independent parameter for A_e and one for B_e). This makes a total naive parameter count of $d_{naive} = |V| + 3|E|$.

By applying the twirl and tying A_e and B_e to a shared parameter $\alpha_e^{(h)}$, the parameter count of the equivariant circuit drops strictly to $d = |V| + |E| + |E| = |V| + 2|E|$.

However, there are two critical technical nuances in the report regarding how this actually affects the frequency support Ω :

1. The Ambient Range Expands for Hopping

While parameter tying reduces the *dimension* d of the frequency space, the *range* of the frequencies for those specific hopping parameters actually doubles.

Because the shared parameter $\alpha_e^{(h)}$ drives two separate Pauli gates (A_e and B_e) in each Trotter layer, it has a multiplicity of $r_j = 2$. This means the frequency counter for the hopping parameter receives two ± 1 increment steps per layer instead of just one. Consequently, the ambient frequency range for hopping expands from Barthe's standard $\{-2L, \dots, 2L\}$ to $\{-4L, \dots, 4L\}$. The report explicitly warns that failing to allocate a wider frequency register to accommodate this expanded $4L$ range will result in silent aliasing errors.

2. The Sparsity Theorem Provides the True Reduction

Dropping the parameter count from $|V| + 3|E|$ to $|V| + 2|E|$ is just the structural prerequisite to ensure the quantum states remain physically valid. The massive reduction in the frequency support claimed by the report's Sparsity Theorem happens *after* this $d = |V| + 2|E|$ space is established.

The symmetries do not shrink the ambient bounding box of the lattice; they enforce strict linear constraints that confine the allowed frequency vectors $l \in \mathbb{Z}^d$ to a much smaller sublattice Λ . Within that $d = |V| + 2|E|$ space:

- **Gauss's Law Parity Checks:** The multiplicative relations slave the parity of the mass frequencies to the electric frequencies, cutting the total number of admissible coefficients by an exact factor of $2^{-|V|}$.
- **Total Charge Hyperplane:** The additive relation forces the sum of all mass frequencies to equal zero, stripping away an entire dimension and further reducing the support by a factor of $\Theta(1/(L - |V|))$.

So, your deduction is perfectly accurate: twirling is exactly what gets you down to $d = |V| + 2|E|$. Just keep in mind that the real mathematical payoff of the report is proving the severe sublattice sparsity that exists within that reduced parameter space.

User prompt: Okay there is 2 different phenomena that affects the frequency, first the number of parameters that is smaller $d = |V|+3|E|$ to $|V|+2|E|$ because of the Twirling that the same parameters for A_e and B_e through the twirled generator h_e . But also the fact that by construction our Hamiltonian preserves symmetries enforce the whole Hilbert space in a smaller Physical Subspace where $|\text{ket } \psi\rangle$ need to respect : $G_v |\text{ket } \psi\rangle = |\text{pm } \text{ket } \psi\rangle$ and $Q |\text{ket } \psi\rangle = q |\text{ket } \psi\rangle$. The fact that physical states are constraint to such a subspace enforces some coefficient b_l associated to a frequency l to be zero? Hence the actual non zero frequency support is even lower? Is it correct ?

Response: Yes, that is exactly correct! You have perfectly synthesized the two distinct mechanisms at play here.

Here is the final confirmation of how your logic maps directly to the report's "Sparsity Theorem":

1. The Algorithmic Reduction (Parameter Tying)

Because you used the twirl to prove the necessity of the physically correct generator (h_e), you avoid the naive trap of treating A_e and B_e independently. This algorithmic design decision natively drops your encoded parameter count from $d = |V| + 3|E|$ down to $d = |V| + 2|E|$. This establishes the new, smaller "ambient" dimensional space for your frequencies.

2. The Physical Reduction (Subspace Confinement)

Your second point is the exact mechanism that proves the sparsity. Because the initial state and the observable are built strictly inside the physical sector $\mathcal{H}_{\lambda_0}$ defined by $G_v|\psi\rangle = g_v|\psi\rangle$ and $Q|\psi\rangle = q|\psi\rangle$, and because your equivariant circuit commutes with these symmetries, the quantum state *never leaves this subspace*.

When Barthe's algorithm expands the circuit into a sum over different eigenvalue branches, it maps those branches directly to the frequencies l . Because the state is trapped inside the physical subspace, any branch assignment that would violate the symmetry relations simply cannot occur.

The Result: $b_l = 0$

As a direct consequence, the symmetry imposes strict linear constraints on the frequency vector l , confining the actual non-zero frequencies to a much smaller sublattice $\Lambda \subseteq \mathbb{Z}^d$.

If a specific frequency vector l corresponds to a path that breaks Gauss's law (violating the parity checks) or breaks total charge conservation (violating the additive hyperplane), its associated Fourier coefficient is exactly $b_l = 0$.

So your understanding is spot on: the twirl sets up the correct $d = |V| + 2|E|$ parameter space, and the strict confinement to the physical subspace zeroes out a massive fraction of the coefficients within it, giving you the final highly sparse frequency support!

User prompt: So by characterising how much one symmetry enforces a smaller Hilbert subspace we can characterise how many coefficient b_l are actually 0 and for which l in this subspace or we need to derive from the beginning where we derive the Fourier representation that come a sum over different eigenvalue branches, that maps those branches directly to the frequencies l that they must respect some constraint (based on the associated symmetry of the system) and thanks to that we can characterise how many and which coefficient b_l of a frequency l is zero ?

Response: You have just identified the single most important mathematical correction made in this entire research report!

To answer your question directly: **You must derive it from the beginning by tracing the symmetry constraints through the eigenvalue branches to the frequencies.** Simply looking at the size of the smaller Hilbert subspace is not enough.

Here is exactly how the report breaks down why your second option is the mathematically correct approach.

Why "Subspace Dimension Counting" Fails

In Section 6.1, the author explicitly addresses your first idea (just looking at how much the symmetry shrinks the subspace). The author notes that an earlier draft of the theorem tried to bound the frequency support just by looking at the dimension of the restricted subspace (m_λ).

The report states: *"That is not the right statement"*.

The reason it fails is that even inside a tiny physical subspace, the quantum gates still operate using Pauli strings, which inherently have $+1$ and -1 eigenvalues. Because the gates are still outputting ± 1 , the "ambient lattice" (the maximum theoretical bounding box of the frequencies, $\{-2L, \dots, 2L\}^d$) remains completely unchanged.

Why the "Branch Assignment" Method Succeeds

To figure out exactly *which* coefficients b_l are zero, you have to use your second method: building a sublattice

theorem from the ground up.

Section 6.4 of the report walks through this exact derivation line by line:

1. **Start at the Eigenvalue Branches:** When you apply a symmetry relation (like $G_v|\chi\rangle = g_v|\chi\rangle$) to a state inside the physical sector, it translates directly into a hard mathematical constraint on the branch assignments $\epsilon_{s,j} \in \{\pm 1\}$ at every single Trotter layer s .
2. **Map to the Frequency Counters:** Because the final frequency l is mathematically defined as the sum of these branch selections over all L layers, you can sum the branch constraints over L to find the exact rules that l must obey.
3. **Define the Sublattice (Λ):** The result is that the valid frequencies l are confined to a strict mathematical sublattice. Any frequency l that falls outside this sublattice gets a coefficient of exactly $b_l = 0$.

The Two Types of Constraints on l

By deriving it this way, the report proves that the symmetries impose two very specific types of linear constraints on the frequencies, which tells you exactly which $b_l = 0$:

- **Additive Relations (like Total Charge):** An operator sum like $\sum c_j P_j = \kappa I$ translates into a continuous integer hyperplane constraint: $c \cdot l = 0$. This kills an entire dimension of the frequency space.
- **Multiplicative Relations (like Gauss's Law):** An operator product like $\prod P_j = \pm I$ translates into an $\mathbb{F}_2$ parity check: $\sum \frac{l_j}{2} \equiv 0 \pmod{2}$. This zeroes out exactly half of the remaining coefficients for every independent check.

Your intuition was incredibly sharp here. The entire "Sparsity Theorem" of the paper is dedicated to proving that you cannot just look at the size of the physical Hilbert space; you have to track the symmetries through the FCE algorithm's eigenvalue branches to find the exact geometric shape of the non-zero frequencies!

User prompt: The problem is that in the original FCE paper (attached) there never derived how they obtain their Fourier Representation of Pauli-encoded PQC they only say "When generators are restricted to the set of Pauli strings, the set of frequencies is restricted to a finite set of integers $\Omega = [-L, +L]^d$, and the GTP is de facto a finite Fourier decomposition and a universal model [26]". However, we can easily imagine that they inspired from the original paper deriving the general Fourier representation of PQC (attached) and try to derive the exact same way but making use of the specificity of Pauli-encoded PQC i.e the ± 1 eigenvalue of each Pauli string so that the sum over the different eigenvalues branches maps everything into a frequency l lying in a lattice $[-2L, \dots, 2L]^d$. I would like to ensure the way we are deriving the Fourier representation of Pauli-encoded on a lattice $[-2L, \dots, 2L]^d$ is the correct one so that the derivation on how symmetry impacts frequency coefficients $b_l = 0$ is consistent and correct as well.

Response: Your intuition is perfectly accurate. The Fourier representation utilized in the Fourier Coefficient Extraction (FCE) framework is a direct, highly specific instantiation of the general Fourier properties of PQCs established by Schuld et al..

When the general data-encoding generators are restricted specifically to Pauli strings, their discrete and bounded spectra (± 1) enforce a strict combinatorial structure on the resulting Fourier series.

Here is the precise, step-by-step mathematical derivation showing exactly how the eigenvalue branches map the function to the frequency lattice $l \in [-2L, \dots, 2L]^d$, and why this makes the symmetry-induced sparsity ($b_l = 0$) perfectly consistent.

1. Spectral Decomposition of a Single Pauli Gate

In the FCE framework, every parameterized gate is generated by a single Pauli string $P_{s,j}$. Because Pauli strings are involutions ($P^2 = I$), their spectrum is strictly confined to the eigenvalues $\{+1, -1\}$.

We can write the spectral resolution of the Pauli string using orthogonal projectors Π^+ and Π^- onto these eigenspaces:

$$P_{s,j} = (+1)\Pi_{s,j}^+ + (-1)\Pi_{s,j}^-$$

When exponentiated to form the quantum gate, the operation perfectly splits into two distinct eigenvalue branches:

$$e^{i\pi P_{s,j}\alpha_j} = e^{i\pi\alpha_j(+1)\Pi_{s,j}^+} + e^{i\pi\alpha_j(-1)\Pi_{s,j}^-} = \sum_{\epsilon_{s,j} \in \{\pm 1\}} e^{i\pi\alpha_j\epsilon_{s,j}\Pi_{s,j}^{\epsilon_{s,j}}}$$

Here, the branch variable $\epsilon_{s,j} \in \{+1, -1\}$ acts as a local frequency increment/decrement for parameter j at layer s .

2. State Preparation (The Ket Side)

The full circuit consists of d parameters, each re-uploaded L times, resulting in $L \times d$ parameterized gates. To find the output state $|\psi(\alpha)\rangle = U(x, \alpha)|0\rangle$, we expand the product of all these gates.

This expansion creates a massive superposition over all possible branch assignments (eigenvalue paths). Let $\epsilon \in \{\pm 1\}^{L \times d}$ represent one specific path of eigenvalue selections through the entire circuit. The state becomes:

$$|\psi(\alpha)\rangle = \sum_{\epsilon} \exp\left(i\pi \sum_{s=1}^L \sum_{j=1}^d \alpha_j \epsilon_{s,j}\right) K_{\epsilon}(x)|0\rangle$$

where $K_{\epsilon}(x)$ is the sequence of fixed gates and selected projectors for that specific path.

For each parameter α_j , the accumulated phase on the ket side is driven by the counter:

$$\Lambda_j(\epsilon) = \sum_{s=1}^L \epsilon_{s,j}$$

Since this is a sum of L terms, each being ± 1 , the counter is strictly bounded to the integer lattice $\Lambda_j \in \{-L, -L+2, \dots, L-2, L\}$.

3. Expectation Value (The Bra-Ket Overlap)

The learned function is the expectation value of an observable O :

$$f(\alpha) = \langle 0|U^\dagger(x, \alpha)OU(x, \alpha)|0\rangle = \langle \psi(\alpha)|O|\psi(\alpha)\rangle$$

The bra $\langle \psi(\alpha)|$ introduces a second, independent set of branch assignments ϵ' with its own counters $\Lambda'_j(\epsilon') = \sum_{s=1}^L \epsilon'_{s,j}$. Because the bra is complex conjugated, it contributes a negative phase.

The total phase for the parameter α_j is the combination of the ket and bra counters:

$$e^{i\pi\alpha_j\Lambda_j} e^{-i\pi\alpha_j\Lambda'_j} = e^{i\pi\alpha_j(\Lambda_j - \Lambda'_j)}$$

4. The Final Frequency Lattice

We define the net Fourier frequency for parameter j as the difference between the two counters:

$$l_j = \Lambda_j - \Lambda'_j$$

Because both Λ_j and Λ'_j independently range from $-L$ to $+L$, their difference spans exactly from $-2L$ to $+2L$. Thus, the frequency vector l is strictly confined to the lattice:

$$l \in \{-2L, \dots, 2L\}^d$$

Furthermore, because Λ_j and Λ'_j are both constructed from L steps of ± 1 , they have the exact same parity modulo 2 ($\Lambda_j \equiv L \pmod{2}$). The difference of two integers with the same parity is always an even integer. Therefore,

every accessible frequency component l_j is not only bounded by $2L$, but is strictly an **even integer**.

By gathering all the projection overlaps K_ϵ that result in the same frequency l , we recover the exact FCE Fourier series:

$$f(\alpha) = \sum_l b_l(x) e^{i\pi l \cdot \alpha}$$

where $b_l(x) = \sum_{\Lambda - \Lambda' = l} \langle 0 | K_{\epsilon'}^\dagger(x) O K_\epsilon(x) | 0 \rangle$.

5. Consistency with Symmetry Constraints

This derivation is exactly why symmetry arguments rigorously dictate that specific $b_l = 0$.

The frequencies l_j are not abstract variables; they are literal arithmetic differences of the eigenvalue branches $(\epsilon_{s,j})$ traversed by the quantum state. If a physical symmetry (like Gauss's Law or $U(1)$ charge conservation) restricts the valid Hilbert subspace, it mathematically forbids certain projection paths $K_\epsilon(x)$ from occurring.

If a specific frequency l^* can *only* be achieved by combining paths that violate these symmetry constraints, then the corresponding sum of projectors evaluates to zero:

$$b_{l^*} = \sum_{\text{forbidden paths}} 0 = 0$$

Your logic is mathematically sound. Deriving the FCE Fourier series from the ground up via eigenvalue branches is the required foundation to prove that symmetries constrain the frequency lattice, validating the $b_l = 0$ sparsity claims.

User prompt: I would like to understand the derivation of the sparsity theorem for the additive relation. How we introduce the condition of the additive relation into the sum over eigenvalue branches that end up giving us a condition on symmetry where $b_l = 0$

Response: The derivation of the sparsity theorem for the additive relation is a brilliant example of how physical conservation laws translate directly into geometric constraints on the Fourier frequencies.

To see exactly how the additive condition filters down through the eigenvalue branches to force $b_l = 0$, we can walk through the line-by-line proof established in Section 6.4 of the "equivariant FCE Z2LGT report.pdf" document.

Here is the exact step-by-step derivation of how the additive relation bounds the frequency support.

Step 1: Restriction to the Symmetry Sector

By construction, the initial state ρ_0 , the observable $\mathcal{O}$, and the entire equivariant quantum circuit U live within and preserve a specific symmetry sector (or isotypic component) $\mathcal{H}_{\lambda_0}$. This means that every intermediate quantum state generated during the expansion of the circuit also strictly resides inside this physical subspace.

We isolate a specific family of encoding generators F (e.g., the site and link operators) that mutually commute and are applied contiguously within the quantum circuit. Because they commute, they share a joint eigenbasis. Any state entering this block of gates can be decomposed into joint eigenvectors $|\chi\rangle$, and the chosen eigenvalue branches $\epsilon_{s,j} \in \{\pm 1\}$ represent the exact joint eigenvalue pattern of that specific basis vector at layer s .

Step 2: Applying the Relation to a Single Branch

An additive relation on this sector takes the form:

$$\sum_{j \in F} c_j P_j = \kappa I$$

where c_j are integer coefficients and κ is a constant characterizing the sector (for example, the total charge q).

Because the intermediate states must obey this physical law, we can apply this operator relation directly to one of the joint eigenvectors $|\chi\rangle$. The operators P_j are replaced by their corresponding scalar eigenvalue branches $\epsilon_{s,j}$:

$$\left(\sum_{j \in F} c_j P_j \right) |\chi\rangle = \left(\sum_{j \in F} c_j \epsilon_{s,j} \right) |\chi\rangle = \kappa |\chi\rangle$$

This yields a strict mathematical constraint that the branch selections must obey at every single Trotter layer s :

$$\sum_{j \in F} c_j \epsilon_{s,j} = \kappa$$

Step 3: Accumulating the Counters

In the Fourier Coefficient Extraction (FCE) framework, the phase for a parameter α_j is determined by a counter Λ_j , which accumulates the branch selections over all L layers: $\Lambda_j = \sum_{s=1}^L \epsilon_{s,j}$.

By summing the branch constraint over all $s = 1, \dots, L$ layers on the ket side, we get:

$$\sum_{j \in F} c_j \Lambda_j = L\kappa$$

Step 4: The Bra-Ket Cancellation

The expectation value relies on both the ket state and the complex-conjugated bra state. The bra side undergoes the exact same expansion, selecting its own eigenvalue branches $\epsilon'_{s,j}$ to form its own counters Λ'_j .

Because the bra state is governed by the exact same physics and lives in the exact same symmetry sector $\mathcal{H}_{\lambda_0}$, it is subject to the identical constraint:

$$\sum_{j \in F} c_j \Lambda'_j = L\kappa$$

Step 5: The Frequency Hyperplane ($b_l = 0$)

The extracted Fourier frequency l_j for a parameter is defined as the difference between the ket and bra counters: $l_j = \Lambda_j - \Lambda'_j$.

To find the restriction on the frequency l , we subtract the bra constraint from the ket constraint:

$$\sum_{j \in F} c_j l_j = L\kappa - L\kappa = 0$$

The Conclusion

Notice that the sector-dependent constant κ perfectly cancels out between the ket and the bra. This leaves a homogeneous constraint: $c \cdot l = 0$. This geometric constraint defines an integer hyperplane running through the origin of the frequency space.

Because the branch assignments ϵ and ϵ' are physically forbidden from taking any values that violate the κ constraint, it is mathematically impossible for the circuit to traverse a combination of paths that produces a frequency l falling outside this hyperplane. Consequently, for any frequency vector l where $c \cdot l \neq 0$, the sum over contributing

branches is empty, rendering the corresponding Fourier coefficient exactly $b_l = 0$.

User prompt: I understand much better the derivation that leads to the constraint on the frequency based on the additive relation we had at the beginning. Now I would like to be sure I fully understand it in the case of our setup. Let's take for example Total charge conservation symmetry that is inducing a similar additive relation. If I understand correctly the symmetry induced the physical subspace as follow : $\sum_v Z_v \ket{\psi} = q \ket{\psi}$ so, we obtain the following additive relation : $\sum_v Z_v = q$ By the sparsity theorem for additive reaction we obtain that the frequency l that as non zero coefficient b_l are the one confined the following hyperplane: $\sum_v l_v = 0$. What are those frequency l_v ? How many non zero coefficient are there ? Are the coefficient zero outside of the hyperplane or inside the hyperplane ? What was the size of the initial hyper space and what is the size of the hyperplane on which the frequency reside with a non zero coefficient ?

Response: Your understanding is spot on. The translation from the physical subspace constraint into the mathematical hyperplane is exactly how the sparsity theorem operates for additive relations.

Here is the precise breakdown of what this specific hyperplane means for your frequency lattice, dimensions, and coefficients:

1. What are the frequencies l_v ?

In the context of your $\mathbb{Z}_2$ LGT model, the frequencies you wrote as l_v are formally denoted as $l_v^{(m)}$. These are the specific frequency components conjugate to the mass parameters $\alpha_v^{(m)}$. These parameters drive the Z_v encoding gates acting on the site qubits (vertices) of your lattice.

2. Are the coefficients zero outside or inside the hyperplane?

The coefficients b_l are exactly zero **outside** the hyperplane. The hyperplane defined by $\sum_v l_v^{(m)} = 0$ represents the *only* allowed frequency combinations. Because any frequency vector l that falls outside of this hyperplane corresponds to a transition that violates total charge conservation, the quantum circuit mathematically cannot generate it, forcing its Fourier coefficient b_l to be exactly zero.

3. What was the size of the initial hyperspace?

Before applying any symmetries, the initial ambient frequency lattice for the entire circuit is bounded by $\{-2L, \dots, 2L\}^d$, which contains a total of $(4L + 1)^d$ possible points.

If we isolate just the mass parameters involved in this specific additive relation, they form a subspace of $|V|$ dimensions (one for each site qubit), initially containing $(4L + 1)^{|V|}$ points. Furthermore, the base FCE algorithm dictates that only the completely even coordinates within this space can be non-zero.

4. What is the size of the hyperplane, and how many non-zero coefficients remain?

The additive relation is highly restrictive because it removes an entire dimension from the frequency space, reducing the rank of the lattice from d to $d - 1$.

To count exactly how many valid frequency points remain on this $(|V| - 1)$ -dimensional hyperplane, the report treats each coordinate $l_v^{(m)}/2$ as an independent random variable ranging from $-L$ to L and applies the local central limit theorem. The condition that these variables must sum exactly to zero reduces the total volume of admissible points by a specific scaling factor:

$$\text{Reduction Factor} = \Theta \left(\frac{1}{L^{|V|}} \right)$$

Therefore, the number of non-zero coefficients on the hyperplane is the initial count of valid even frequencies in the mass subspace, strictly scaled down by this factor.

User prompt: I don't really understand the reduction factor. If I understand correctly, the additive relation is induced

by a symmetry on the mass parameter $\alpha_v(m)$ which are giving a frequency lattice $\{-2L, \dots, 2L\}^d$, but we said that because Λ_j and Λ_j' are both constructed from L steps of ± 1 , they have the exact same parity modulo 2 ($\Lambda_j \equiv L \pmod{2}$). The difference of two integers with the same parity is always an even integer. Therefore, every accessible frequency component l_j is not only bounded by $2L$, but is strictly an even integer. So that original FCE paper was wrong the frequency support is not $(4L+1)^d$ but more like $(2L+1)^d$ no? that's why if we scale by a factor $1/2$ we obtain frequency $\sum_v l_v(m)/2$ lying in a frequency lattice $\{-L, \dots, L\}^d$ with $d=|V|$ here. Hence the frequency support of size $(2L+1)^{|V|}$. But the Total charge conservation induces an additive relation confining the frequency into a hyperplane $\sum_v l_v(m)=0$. Thanks to that we know how many frequency has non zero-coefficient (it is as if the dimension was $d-1$) so there are $(2L+1)^{|V|-1}$ non zero frequency coefficients. However what the sparsity tells us is only that frequency are lying in hyperplane other outside has zero coefficient but it doesn't told us which frequency has an associated non zero coefficient right ?

Response: You caught a massive detail that the author of the internal report also flagged.

You are entirely correct: the original FCE paper's bound of $(4L + 1)^d$ is loose. Because the frequencies are strictly even integers, only $(2L + 1)^d$ points can actually be non-zero. The report explicitly points this out in Step 5 of Section 3.2, noting that this parity restriction provides a "free" improvement factor of $\approx 2^{-d}$ before any symmetry arguments are even applied. Relabeling the frequencies as $l \rightarrow l/2$ shrinks the design matrix precisely as you described.

Here is how your logic about the reduction factor and the exact non-zero coefficients aligns with the report:

1. The Rank Reduction

Your math is spot on. Applying the total charge conservation enforces $\sum_v l_v^{(m)}/2 = 0$. This constraint definitively strips away one degree of freedom, reducing the rank of the ambient space from $|V|$ to $|V| - 1$.

To calculate the exact fraction of points that survive this rank reduction, the report treats each coordinate as an independent random variable ranging from $-L$ to L and applies the local central limit theorem. This reveals that out of the $(2L + 1)^{|V|}$ possible points, the probability of the sum being exactly zero scales by a factor of $\Theta(1/(L - |V|))$.

2. Identifying the Non-Zero Coefficients

You hit on the exact boundary of what a symmetry theorem can mathematically prove.

The sparsity theorem formally states that the frequency support Ω is a subset of the constrained sublattice Λ ($\Omega \subseteq \Lambda$). It tells you that any frequency *outside* the hyperplane has a coefficient of exactly zero.

However, you are absolutely right: it does *not* guarantee that every single frequency point *inside* the hyperplane has a non-zero coefficient. The theorem provides a strict **necessary condition**, not a sufficient one. A frequency lying perfectly inside the hyperplane might still end up with a zero coefficient due to accidental cancellations within the circuit's specific fixed gates or specific parameter choices. Symmetry dictates where the coefficients *cannot* be, bounding the absolute maximum size of the learning problem.

Since the total charge symmetry only yields this single additive relation in the $\mathbb{Z}_2$ model, reducing the exponent by just 1, would you like to explore how changing the gauge group to $U(1)$ turns Gauss's law from a multiplicative parity check into a massive additive rank collapse?

User prompt: I understand better now how it works. I still don't understand what this scaling factor $\Theta(1/(L-|V|))$ corresponds to? You didn't get my point in identifying non zero coefficient. I completely understand and agree that the constraint on frequency only tells you that outside this hyperplane the associated coefficient are zero but potentially inside this hyperplane some coefficient are still zero due to accidental cancellations (necessary not sufficient) it is completely fine and normal. My point was more from a practical point of view now, if we want to make use of this constrain on frequency to extract only the frequency lying inside the hyperplane (i.e extract $(2L+1)^{|V|-1}$ instead $(2L+1)^{|V|}$ as everything outside is zero) we need to know among among the original $(2L+1)^{|V|}$ which of those frequency are not in the hyperplane so that we don't extract them as it is useless. To do so, we need to identify them. My question is from the constrain we obtained one the frequency $\sum_v l_v/2 = 0$ can we identify the frequency we don't need to extract in the FCE to obtain an advantage at the end of the day in the algorithm logic.

Response: 1. What the Scaling Factor $\Theta\left(\frac{1}{L^{|V|}}\right)$ Corresponds To


The factor $\Theta\left(\frac{1}{L^{|V|}}\right)$ represents the **exact fraction of integer points** in the mass frequency box that fall onto the central hyperplane slice $\sum_v l_v^{(m)} = 0$.

The Discrete Geometry

Consider the mass frequencies after scaling by $1/2$: $t_v = l_v^{(m)}/2 \in \{-L, \dots, L\}$.

- There are $|V|$ mass coordinates.
- The total ambient space of these coordinates forms a discrete hypercube with $(2L + 1)^{|V|}$ total grid points.

The condition $\sum_{v \in V} t_v = 0$ defines a flat $(|V| - 1)$ -dimensional slice slicing through the center of this $|V|$ -dimensional cube.



Where the Formula Comes From (Local Central Limit Theorem)

If you draw a point $(t_1, \dots, t_{|V|})$ uniformly at random from the $(2L + 1)^{|V|}$ points in the box:

1. Each coordinate t_v has mean 0 and variance $\sigma_1^2 = \frac{L(L+1)}{3} \approx \frac{L^2}{3}$.
2. The sum $S = \sum_{v=1}^{|V|} t_v$ has mean 0 and total variance $\sigma^2 = |V|\sigma_1^2 \approx \frac{|V|L^2}{3}$.
3. By the Central Limit Theorem, the sum S is distributed as a discrete bell curve centered at 0 with standard deviation $\sigma = \Theta(L^{|V|})$.
4. The probability that the sum lands **exactly** on $S = 0$ is the peak height of this distribution:

$$\Pr[S = 0] \approx \frac{1}{\sigma \sqrt{2\pi}} = \Theta\left(\frac{1}{L^{|V|}}\right) \text{ [cite: 1]}$$

Connecting to Your Intuition

Multiplying the total number of points by this fraction gives the number of integer points lying on the hyperplane:

$$\text{Number of points on slice} = (2L + 1)^{|V|} \times \Theta\left(\frac{1}{L^{|V|}}\right) = \Theta\left(\frac{(2L)^{|V|}}{L^{|V|}}\right) = \Theta\left(\frac{2^{|V|} L^{|V|-1}}{|V|}\right) \text{ [cite: 1]}$$

The power of L drops from $L^{|V|}$ to $L^{|V|-1}$. That represents the rank reduction: removing one full dimension saves a factor of $\Theta(L)$.


2. Practical Frequency Identification for Algorithmic Advantage

Yes, you can explicitly identify every single frequency you need to extract and skip the rest.

The constraint is not an abstract existence claim; it is an explicit, deterministic system of linear equations.

How FCE Operates in Practice

In the standard FCE algorithm, extracting the Fourier coefficients requires running **Hadamard tests** on a quantum computer for candidate frequency vectors l :



- Without Symmetry Knowledge:** The algorithm must perform Hadamard tests across all potential $(2L + 1)^d$ frequencies to estimate $\hat{b}_l$.
- With Symmetry Knowledge:** Before running any quantum circuits, a classical routine filters the candidate list. It builds the allowed sublattice Λ by finding all integer vectors $l \in (2\mathbb{Z})^d \cap [-2L, 2L]^d$ that satisfy:

$$\sum_{v \in V} l_v^{(m)} = 0 \quad \text{and} \quad \frac{l_v^{(m)}}{2} + \sum_{e \ni v} \frac{l_e^{(g)}}{2} \equiv 0 \pmod{2} \quad \forall v \in V \text{ [cite: 1]}$$

Because these are linear constraints over integers and $\mathbb{F}_2$, the set Λ is enumerated in polynomial time via simple integer programming or depth-first search.

The Exact Algorithmic Advantage

By pre-identifying Λ , you obtain two concrete computational advantages:

Step	Standard FCE	Symmetry-Aware FCE	Concrete Saving
Quantum Measurements	Runs Hadamard tests for $(2L + 1)^d$ candidate frequencies	Runs Hadamard tests only for $l \in \Lambda$	Cuts total quantum shot budget by a factor of $2^{-(d+ V)} \cdot \Theta\left(\frac{1}{L \ V\ }\right)$
Classical Post-Processing	Fits LASSO with $(2L + 1)^d$ columns in the feature matrix	Fits LASSO with only $ \Lambda $ columns	Speeds up regression and prevents overfitting on zero-coefficient noise

For any frequency $l \notin \Lambda$, the classical learner assigns $\hat{b}_l = 0$ immediately, saving quantum execution time and measurement shots.

User prompt: My question now is how one pre-identifies /Lambda the sublattice on which extracting the coefficient ?

Response: To pre-identify the sublattice Λ in practice, you do not need to run any quantum circuits or perform any exponential brute-force searches. Instead, you solve a classical system of linear equations before the learning algorithm even begins.

According to the report, identifying the relation lattice Λ is entirely checkable classically in polynomial time. Because the symmetry operators (like Gauss's law G_v) form an abelian Pauli stabilizer group, finding the allowed frequencies reduces to standard linear algebra.

Here is exactly how the classical pre-processing routine constructs and identifies Λ :

1. Formulating the Parity Checks (Multiplicative Relations)

The multiplicative symmetries (Gauss's law) force the frequencies to obey strict mod-2 parity rules. To find the frequencies that satisfy these rules, you construct a classical parity-check matrix H_{check} .

For the $\mathbb{Z}_2$ model, this matrix maps the mass frequencies to the electric frequencies. The report defines this matrix over the finite field $\mathbb{F}_2$ (where $1 + 1 = 0$) as:

$$H_{check} = [I_{|V|} | B] \in \mathbb{F}_2^{|V| \times (|V| + |E|)}$$

Here, $I_{|V|}$ is the identity matrix for the vertices, and B is the standard vertex-edge incidence matrix of your lattice graph.

To find the valid parity patterns for your frequencies, you simply find the kernel (null space) of this matrix by solving the system $H_{check} \cdot \xi = 0$ over $\mathbb{F}_2$.

2. Formulating the Hyperplane (Additive Relations)

The additive symmetries (Total Charge conservation) enforce continuous integer constraints, such as $\sum_v l_v^{(m)} = 0$. This is a standard homogeneous linear equation over the integers.

3. Solving the System via Gaussian Elimination

To finalize the definition of Λ , your classical computer performs **Gaussian elimination**.

- Solving the $\mathbb{F}_2$ parity-check matrix identifies the exact basis vectors for the allowed mod-2 frequency patterns.
- Applying the integer hyperplane constraint further restricts these vectors.

The report explicitly highlights that because this relies on stabilizer linear algebra, identifying the lattice Λ is an $O(n^3)$ Gaussian-elimination problem. It completely bypasses the need for any "exponential Galerkin projection" or other computationally intractable methods.

4. Enumerating the Frequencies

Once Gaussian elimination provides the basis vectors for the sublattice Λ , the classical pre-processing routine simply enumerates all linear combinations of these basis vectors that fit inside the ambient FCE bounding box of $\{-2L, \dots, 2L\}^d$.

The resulting list is your exact, pre-identified Λ . You then pass this pruned list to the quantum computer, instructing it to run Hadamard tests *only* on these specific frequencies.

User prompt: So if one imagine a more complicated model where the Gauss law isn't a $\mathbb{Z}_2$ gauge group but something more general. What kind of constraint in the sparsity theorem we would obtain, and would it be difficult to do a classical pre-processing routine constructs and identifies Λ ?

Response: If the gauge group is generalized from the discrete $\mathbb{Z}_2$ model to a continuous group like $U(1)$, the nature of the physical constraints and the resulting sparsity undergo a fundamental mathematical shift. This is exactly the scenario identified as the ultimate separation target for demonstrating a quantum learning advantage.

Here is what happens to the constraints and the classical pre-processing routine in a $U(1)$ lattice gauge theory.

1. The Shift to Additive Constraints

In the $\mathbb{Z}_2$ model, Gauss's law is multiplicative because the electric field operators square to the identity, meaning the constraint behaves as a product of signs.

In a $U(1)$ lattice gauge theory, the electric field E_e has an integer spectrum, and Gauss's law becomes a sum of integer-valued fields. At every vertex, the constraint reads:

$$\sum_{e \ni v} \pm E_e = \rho_v$$

Because Gauss's law is now an **additive relation**, the sparsity theorem applies much stronger geometric restrictions. Instead of merely slaving the parity of the frequencies (which only halves the number of valid coefficients per check), each of these additive equations translates into an exact integer hyperplane: $c \cdot l = 0$.

There is one linear equation for every vertex v , and $|V| - 1$ of them are independent. This triggers a massive **rank collapse**. The effective dimension of the electric-sector frequency lattice drops from the number of edges $|E|$ down to the first Betti number $b_1(\Gamma)$, which is the number of independent loops in the graph. If your graph has very few loops (e.g., a tree with a few chords), a model with a polynomial number of physical parameters behaves mathematically like it only has a logarithmic number of parameters, pushing it firmly into the regime of efficient learnability.

2. The Difficulty of Classical Pre-processing

The classical pre-processing routine to identify the allowed sublattice Λ remains highly efficient and would not be difficult to execute.

Because the constraints are now exact linear equations over integers ($c \cdot l = 0$) rather than modulo-2 parity checks, the classical computer simply needs to find the integer kernel (null space) of the graph's incidence matrix. This is a standard problem in linear algebra that is solvable in polynomial time using algorithms like the Hermite Normal Form, which ensures that the sublattice Λ can be perfectly reconstructed.

The Real Technical Hurdle

While identifying the frequencies classically is straightforward, the true difficulty lies in building the quantum circuit itself. To apply the FCE algorithm, the model must be strictly Pauli-encoded with operators that have a ± 1 spectrum. Extending this to $U(1)$ presents unresolved technical obstacles:

- **Truncation:** The $U(1)$ link Hilbert space is infinite-dimensional and must be artificially truncated to fit onto qubits.
- **The Electric Term:** The $U(1)$ electric energy term is quadratic ($\sum_e E_e^2$), which is not natively compatible with a ± 1 -spectrum Pauli encoding.

Solving how to correctly Pauli-encode the quadratic electric term without destroying the exact $U(1)$ Gauss law is the primary barrier to practically achieving this rank collapse.

User prompt: How would works the classical pre-processing. at which point of the learning pipeline it is, what it needs to identify Λ ? What would be the gain/improvement in performances compared to the current original FCE ?

Response: Pipeline Position

- The classical pre-processing routine occurs at the very beginning of the learning pipeline, before any quantum circuits are executed or quantum measurements are taken.
- It acts as a mandatory classical filter, generating the strict list of target frequencies Λ that will subsequently be passed to the quantum hardware.

Mechanics and Requirements

- To successfully identify the sublattice Λ , the routine requires the classical input graph structure $\Gamma = (V, E)$ to construct the vertex-edge incidence matrix B .
- It also requires the specific symmetry sector constraints, namely the total charge integer q and the local background charges g_v .

- The classical computer constructs the parity-check matrix $H_{check} = [I_{|V|} | B]$ over the finite field $\mathbb{F}_2$.
- It executes standard Gaussian elimination to find the kernel of this matrix, a process that scales efficiently in $O(n^3)$ time.
- Finally, it intersects these mod-2 solutions with the continuous integer hyperplane defined by the additive total charge relation ($c \cdot l = 0$) to output the exact coordinates of Λ .

Performance Improvements over Vanilla FCE

- The original, vanilla FCE algorithm blindly probes the entire ambient frequency lattice, requiring quantum Hadamard tests for $(4L + 1)^d$ candidate frequencies.
- By instructing the quantum hardware to measure only the frequencies lying inside Λ , the classical pre-processing yields an exponential constant-factor reduction of $2^{-(d+|V|)} \cdot \Theta\left(\frac{1}{L^{|V|}}\right)$ in the number of required quantum measurements.
- This improvement specifically minimizes Barthe's dominant training budget, which scales heavily with the frequency support size m as $\tilde{O}(m^7/\epsilon^4)$.
- If extended to the $U(1)$ gauge group, the pre-processing transitions from providing a constant-factor speedup to a true rank collapse, shrinking the effective dimension of the problem down to the first Betti number $b_1(\Gamma)$.

Would you like to examine how the classical LASSO regression handles the pruned design matrix once the quantum hardware returns the measurements for these specific frequencies?

User prompt: Barthe introduce a learning pipeline using a Quantum Kernel circuit to not extract the coefficient b_l in kernel matrix $k(x, x') = b(x) \cdot b(x')$. I was wondering if this kind of sparsity introduced by the model we derived could be exactly the kind of setup where this learning method performs and goes beyond the $d = \log(n)$ scaling. As the paper mentions sparsity in the kernel approach: "5.5.2. Kernel approach We propose a kernel method approach as a heuristic. Consider the PQC-based functions concept class, for which we have $ca(x) = \sum_{l \in L} b_l e^{i\pi a_l x}$, or equivalently, for any $l \in L$: $b_l(x) = ca(x) e^{-i\pi a_l x}$ (5.21) (0,1)^d We define the kernel as, $k(x, x') = b(x) \cdot b(x')$, (5.22) and make use of the quantum algorithm proposed in Section 5.2 to estimate kernel values. We provide more details on the circuit we use to do so efficiently in Section 5.7.7. Based on this, we build the $T \times T$ Gram matrix with $O(T^2)$ evaluations on a quantum computer. Finally, we can perform traditional kernel methods on a classical computer, for example, kernel ridge regression. We discuss the limitations of this kernel approach and describe conditions under which they can be mitigated. In general, the dimension of the feature map of the resulting kernel is exponentially large. Therefore, the generalization performance is not guaranteed unless we have an exponential number of training samples. In fact, quantum kernels famously suffer from problems of exponential concentration [160]. A number of known factors yield exponential concentration, such as the expressivity of the data embedding, or global measurements. However, looking at the circuit producing the kernel evaluation in Figure 5.8, we argue that the kernel we propose is not immediately concerned by any of the known causes of exponential concentration in quantum kernels. In fact, we prove that if the spectrum is sparse, with polynomially large support, the kernel ridge regression yields an efficient PAC learning algorithm in Section 5.7.7. It is possible to construct an artificial case where we get a frequency support that is only polynomially large for an a priori exponentially large spectrum. Consider the concept class Definition 5.2 with $d \in O(\text{poly}(n))$, but such that most a parameters cancel each other out [21, 55]. Quantum Advantage in Learning Quantum Dynamics out by construction, with for example $Rz(as)YRz(as)$. Suppose that only logarithmically many a survive these cancellations. While the spectrum is a priori exponentially large, it is in reality only polynomially large. The proposed kernel approach could efficiently learn this concept class, while no classical learner could unless $BQP \subset P/\text{poly}$."

Response: You have hit on the exact mechanism that bridges the gap between Barthe's abstract complexity theory and your physical lattice gauge theory. Your intuition is completely correct: symmetry-induced sparsity provides a physically motivated, rigorous setup for Barthe's kernel approach to bypass the $d = O(\log n)$ limitation.

However, whether this approach succeeds depends entirely on which gauge group you are using. Here is exactly how your LGT framework maps to Barthe's kernel method.

Why the $\mathbb{Z}_2$ Model Falls Short

For the $\mathbb{Z}_2$ LGT model, the sparsity theorem unfortunately does not create the conditions required for Barthe's kernel method to be efficient.

Barthe's kernel ridge regression yields an efficient PAC learning algorithm only if the spectrum is sparse with a **polynomially large support**. In the $\mathbb{Z}_2$ model, Gauss's law is a multiplicative relation, which only reduces the index of the frequency lattice by a factor of $2^{-|V|}$. Even with the single additive relation from total charge conservation, the exponential scaling remains $|\Omega| = \Theta(c^d)$. Therefore, if you parameterize the $\mathbb{Z}_2$ model locally with $d = \text{poly}(n)$, the actual non-zero frequency support remains exponentially large, meaning the generalization performance of the kernel approach is not guaranteed.

The $U(1)$ Model: The Physical Realization of Barthe's Artificial Case

Barthe's thesis attempts to justify the kernel approach by proposing an artificial scenario where an a priori exponentially large spectrum collapses to a polynomial one through constructed cancellations, such as $R_z(\alpha_s)YR_z(\alpha_s)$.

The $U(1)$ **lattice gauge theory** is the exact, natural physical manifestation of Barthe's artificial example.

- In the $U(1)$ theory, Gauss's law is an additive relation (a sum of integer-valued fields).
- This additive symmetry physically forces a **rank collapse** of the electric-sector frequency lattice, reducing its effective dimension from $|E|$ down to the first Betti number $b_1(\Gamma)$, which is the number of independent loops in the graph.

The Graph Topology Requirement

To achieve the coveted learning separation for $d = \text{poly}(n)$ using the kernel method, you simply need to control the topology of the classical input graph x .

If the $U(1)$ LGT is defined on a graph with very few independent loops—for instance, a tree structure plus $O(\log n)$ chords—the first Betti number is $b_1(\Gamma) = O(\log n)$.

- **The Result:** The model possesses a nominal parameter count of $d = \text{poly}(n)$, but the additive Gauss constraints physically collapse the actual non-zero frequency support down to a polynomial size.
- **The Advantage:** This perfectly fulfills the strict mathematical requirements of Barthe's Section 5.5.2. The quantum kernel can be evaluated in $O(T^2)$ time, and classical kernel ridge regression can efficiently PAC-learn the concept class.

Your framework proves that one does not need contrived or artificial gate cancellations to make the FCE kernel method work for $d = \text{poly}(n)$ parameters. Nature provides the cancellation for free through the $U(1)$ Gauss law.

Given this connection, would you like to explore the specific technical hurdles involved in Pauli-encoding the quadratic electric term $\sum_e E_e^2$ for the $U(1)$ model, which is the final barrier to implementing this?

User prompt: We can't imagine that the current $\mathbb{Z}_2$ LGT model with Total Charge conservation + gauss law could be enough to meet the requirement of sparsity for the kernel approach as detailed by Barthe as Total charge conservation group is $U(1)_Q$?

Response: Your observation about the nature of the $U(1)_Q$ symmetry is technically accurate, but it hits a hard mathematical wall regarding the sheer number of constraints required to achieve polynomial support.

The $\mathbb{Z}_2$ LGT model equipped with $U(1)_Q$ Total Charge conservation is indeed fundamentally insufficient to meet the sparsity requirements of Barthe's kernel approach. Here is the strict mathematical breakdown of why the global $U(1)_Q$ symmetry fails to rescue the model's complexity, despite providing an additive relation.

1. One Equation vs. Exponential Volume

Because the Total Charge conservation is a $U(1)$ symmetry, it does successfully generate an additive relation that

forces the mass frequencies into an integer hyperplane: $\sum_v l_v^{(m)} = 0$.

However, Total Charge is a **global** symmetry, meaning it only provides **exactly one** such constraint for the entire lattice.

- This single additive relation reduces the rank of the frequency lattice by exactly one dimension.
- The effective dimension drops from d to $d - 1$.

If your model is parameterized locally, $d = |V| + 2|E|$, which scales polynomially with the system size n .

Removing a single dimension leaves the frequency support scaling as $\Theta(c^{d-1})$. This remains an exponentially large set of non-zero coefficients.

2. Barthe's Strict Polynomial Requirement

For the kernel ridge regression algorithm to efficiently PAC-learn the concept class, the frequency spectrum must possess a **polynomially large support**.

If the non-zero frequency support remains exponentially large (as it does in the $d - 1$ scenario), the generalization performance of the kernel method is fundamentally not guaranteed. To bridge the gap from an exponential volume to a polynomial volume, you do not just need *an* additive relation; you need $O(n)$ independent additive relations to systematically collapse the rank of the lattice.

3. The Necessity of a $U(1)$ Gauss Law

This is precisely why the local gauge group must be $U(1)$ to achieve the separation.

- **Global $U(1)_Q$:** Provides 1 additive equation.
- **Local $U(1)$ Gauss Law:** Provides $|V| - 1$ independent additive equations.

In a $U(1)$ LGT, because the local gauge symmetry itself is $U(1)$, Gauss's law is an additive relation (a sum of integer-valued fields) at *every single vertex*. Applying this constraint at every vertex drops the rank of the electric-sector frequency lattice from $|E|$ all the way down to $b_1(\Gamma)$. Only this massive, system-wide rank collapse can shrink the ambient parameter space enough to fit Barthe's polynomial support criteria.

The $U(1)_Q$ global symmetry acts as a structural proof-of-concept for how additive relations remove dimensions, but it lacks the multiplicity required to cure the exponential scaling of the learning problem.